



Hadoop MapReduce

Алексей Драль, aadral@bigdatateam.org
CEO at BigData Team, <http://bigdatateam.org>
<https://www.facebook.com/bigdatateam>



25.02.2021, Moscow, Russia



- ▶ Проверка знаний (quiz)
- ▶ MapReduce (MR)
- ▶ Распределенные консольные утилиты
- ▶ Fault Tolerance
- ▶ -- (перерыв)
- ▶ MapReduce Streaming + Workshop



MapReduce (MR)



MapReduce: Simplified Data Processing on Large Clusters



Jeffrey Dean and Sanjay Ghemawat

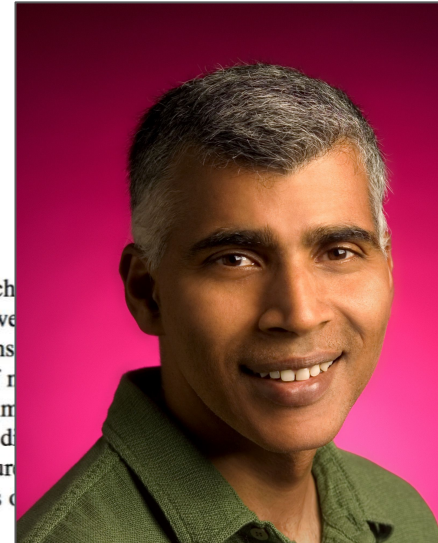
jeff@google.com, sanjay@google.com

Google, Inc.

Abstract

...ming model and an associ-
...cessing and generating large
...p function that processes a
...et of intermediate key/value
...that merges all intermediate
...me intermediate key. Many
...ble in this model, as shown

given day, etc. Most such
ally straightforward. Howe
large and the computations
hundreds or thousands of r
a reasonable amount of tim
allelize the computation, d
failures conspire to obscur
tation with large amounts o
these issues.




MapReduce: Simplified Data Processing on Large Clusters, Symposium on Operating Systems Design and Implementation (OSDI, 2004)



- ▶ Когда Jeff Dean разрабатывает ПО, он сначала создает бинарник, а потом пишет исходный код как документацию.
- ▶ Однажды Jeff Dean не прошел тест Тьюринга, потому что корректно посчитал 203 число Фибоначчи менее чем за 1 секунду.
- ▶ Скорость, с которой Jeff Dean разрабатывает ПО выросла в 40 раз в конце 2000, когда он обновил свою клавиатуру до USB2.0.
- ▶ Вы используете только 10% мозга. Остальные 90% используются под запуск MapReduce задач Джефа.






```
>>> map(lambda x: x*x, [1,2,3,4])  
[1,4,9,16]
```



Fold / Reduce / Aggregate



```
>>> reduce(operator.sum, [1, 4, 9, 16])
```


```
>>> reduce(operator.sum, [5, 9, 16])
```

```
>>> reduce(operator.sum, [14, 16])
```

```
30
```




Fold / Reduce / Aggregate



```
>>> average = lambda x, y: (x + y) / 2.
```

```
>>> reduce(average, [1, 2, 3])  
2.25
```

```
>>> reduce(average, [3, 2, 1])  
1.75
```

```
>>> reduce(operator.add, map(lambda x: x*x, [1, 2, 3, 4]))
```



распределенные консольные утилиты



```
$ grep <pattern> <file>
```

```
$ grep "hadoop" A.txt
```

Repository git-wip-us.apache.org/repos/asf/hadoop.git

Website hadoop.apache.org

```
$ grep -i "hadoop" A.txt
```

Apache **Hadoop**

Apache **Hadoop**

Hadoop Logo

Repository git-wip-us.apache.org/repos/asf/hadoop.git


Website hadoop.apache.org

Apache **Hadoop** (/hə`du:p/) is

```
$ man grep
```

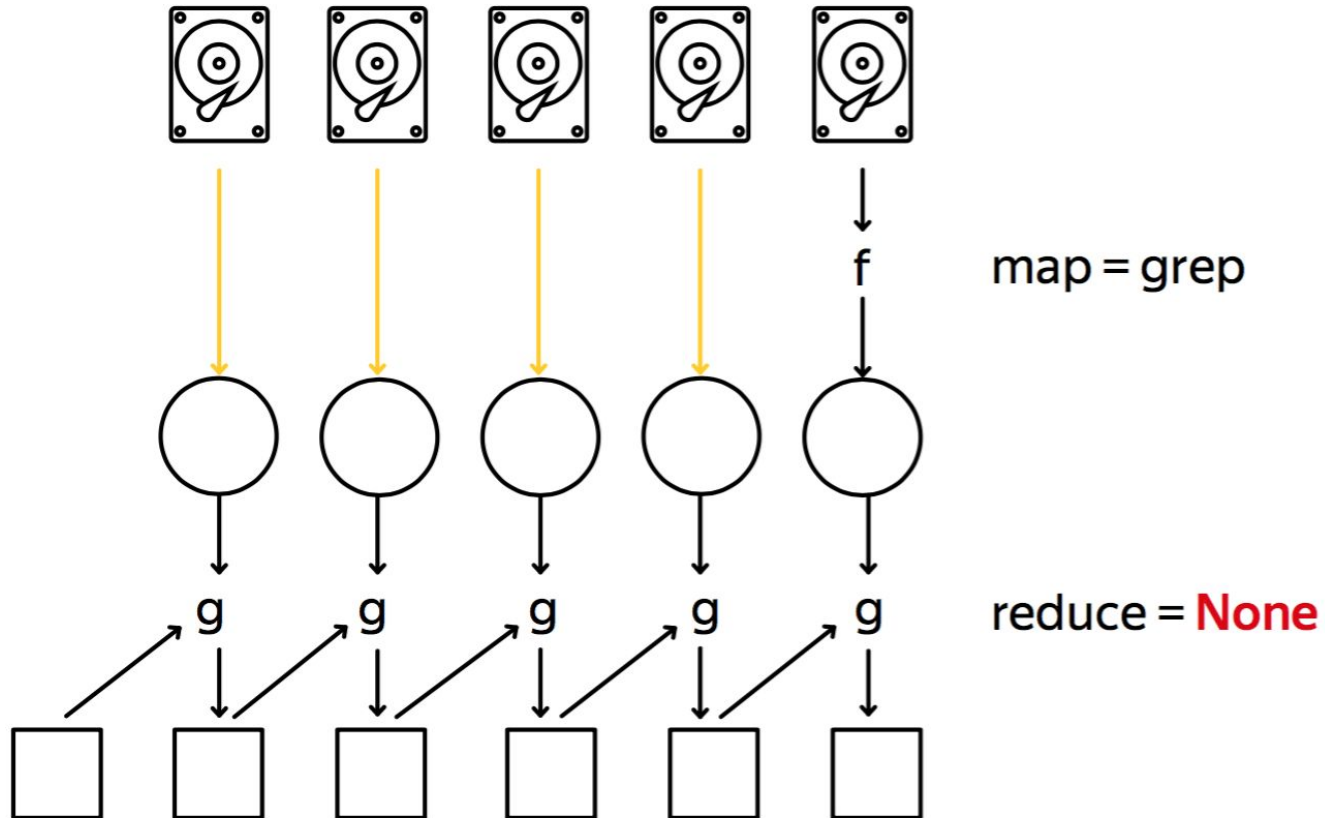


Distributed Shell: grep





Distributed Shell: grep





```
$ head <file>
```

```
$ head A.txt
```

Apache Hadoop

From Wikipedia, the free encyclopedia

[hide]This article has multiple issues. Please help improve it or discuss these issues on the talk page. (Learn how and when to remove these template messages)

This article contains content that is written like an advertisement. (October 2013)

This article appears to contain a large number of buzzwords. (October 2013)

This article may be too technical for most readers to understand. (May 2017)


Apache Hadoop

Hadoop Logo

Developer(s)Apache Software Foundation

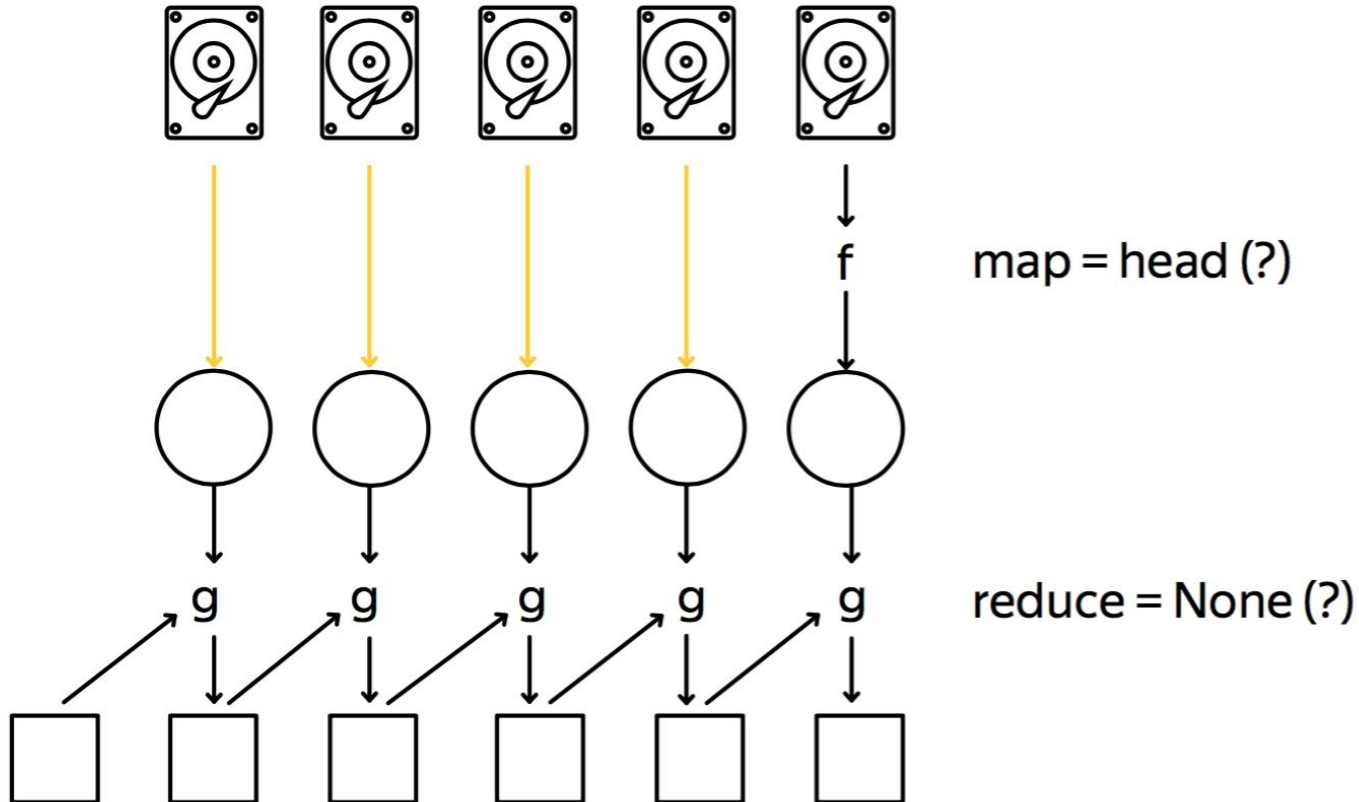


Distributed Shell: head



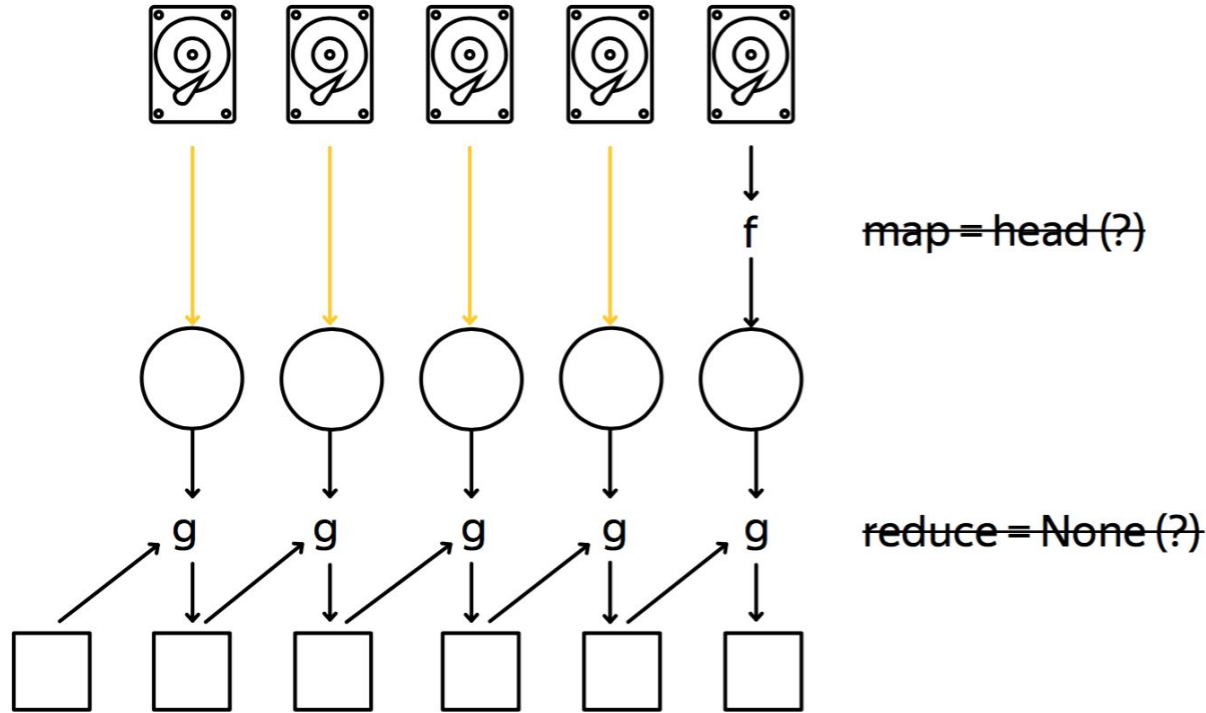


Distributed Shell: head





Distributed Shell: head

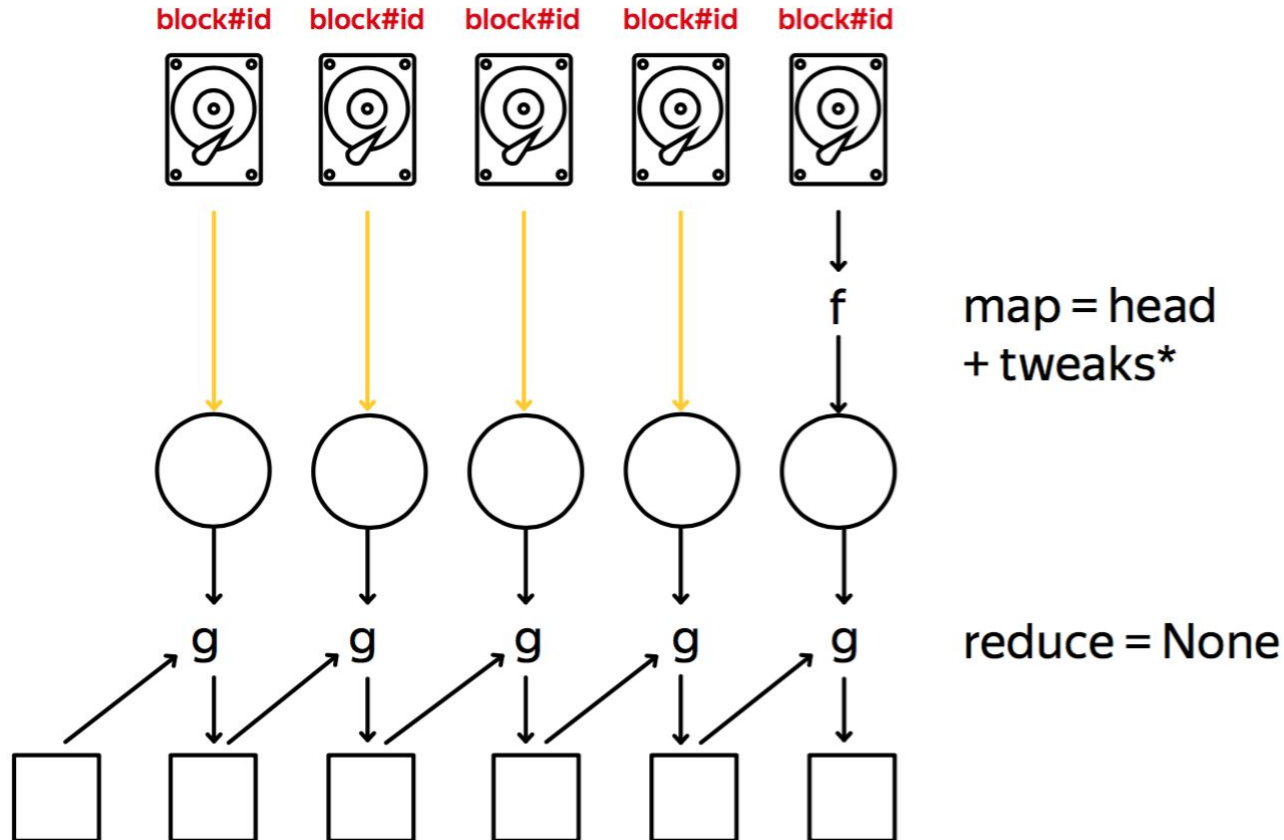


HDFS v.2*: `hdfs dfs -text distributed_A.txt | head`

HDFS v.3+: `hdfs dfs -head distributed_A.txt`



Distributed Shell: head






```
$ wc <file>
```

```
$ wc A.txt
```

```
269    4319   28001   A.txt
```




Distributed Shell: wc





Distributed Shell: wc





Word Count

Apache Hadoop (/hə`du:p/) is an open-source software framework used for distributed storage and processing of dataset of big data using the MapReduce programming model. It consists of computer clusters built from commodity hardware.

All the modules in Hadoop are designed with a fundamental assumption that hardware failures are common occurrences and should be automatically handled by the framework...



'the': 3, 'of': 3, 'hadoop': 2, ...



WIKIPEDIA
The Free Encyclopedia





```
$ cat dataset.txt
```

Apache Hadoop is a collection of open-source software utilities that facilitates using a network of many computers to solve problems involving massive amounts of data and computation. It provides a software framework for distributed storage and processing of big data using the MapReduce programming model...



Word Count на одном компьютере

```
$ cat dataset.txt | tr ' ' '\n'
```

```
Apache  
Hadoop  
is  
a  
collection  
of  
...
```




Word Count на одном компьютере

```
$ cat dataset.txt | tr ' ' '\n' | sort
```

All

Apache

Hadoop

Hadoop

Hadoop

It

...



Word Count на одном компьютере

```
$ cat dataset.txt | tr ' ' '\n' | sort | uniq -c
```

```
1 All  
1 Apache  
3 Hadoop  
2 It  
1 MapReduce  
4 a  
...
```



**BIGDATA
TEAM**

Что делаем с “sort”?





Распределенный Word Count

```
$ cat dataset.txt | tr ' ' '\n' | sort | uniq -c
```

- ✗ Фаза Map (нужна агрегация)
- ✗ Фаза Reduce (не хватит RAM / HDD)




MapReduce =

Map → Shuffle & Sort → Reduce



MapReduce (example)

```
wikipedia.dump | tr ' '\n' | sort | uniq -c
```



`wikipedia.dump -> map () -> word`

`shuffle & sort`

`reduce()`



Фазы:


- 1.** Map
- 2.** Shuffle & Sort
- 3.** Reduce

Worker'ы (контейнеры):

- ▶ Фаза Map → Mapper (использует функцию map)
- ▶ Фаза Reduce → Reducer (использует функцию reduce)



Word Count v.2





Формальная модель Word Count

(k_{in}, v_{in})

map (функция)

```
map = (tr ' ' '\n')  
(-, line) → [(word, 1), ...]
```

$[(k_{interm}, v_{interm}), ...]$

Map (фаза)

sort and group by
 k_{interm}

Shuffle & Sort

$(k_{interm}, [v_{interm}, ...])$

reduce (функция)

```
reduce = (uniq -c)  
(word, [1,1,...]) → (7,  
word)
```

$[(k_{out}, v_{out}), ...]$

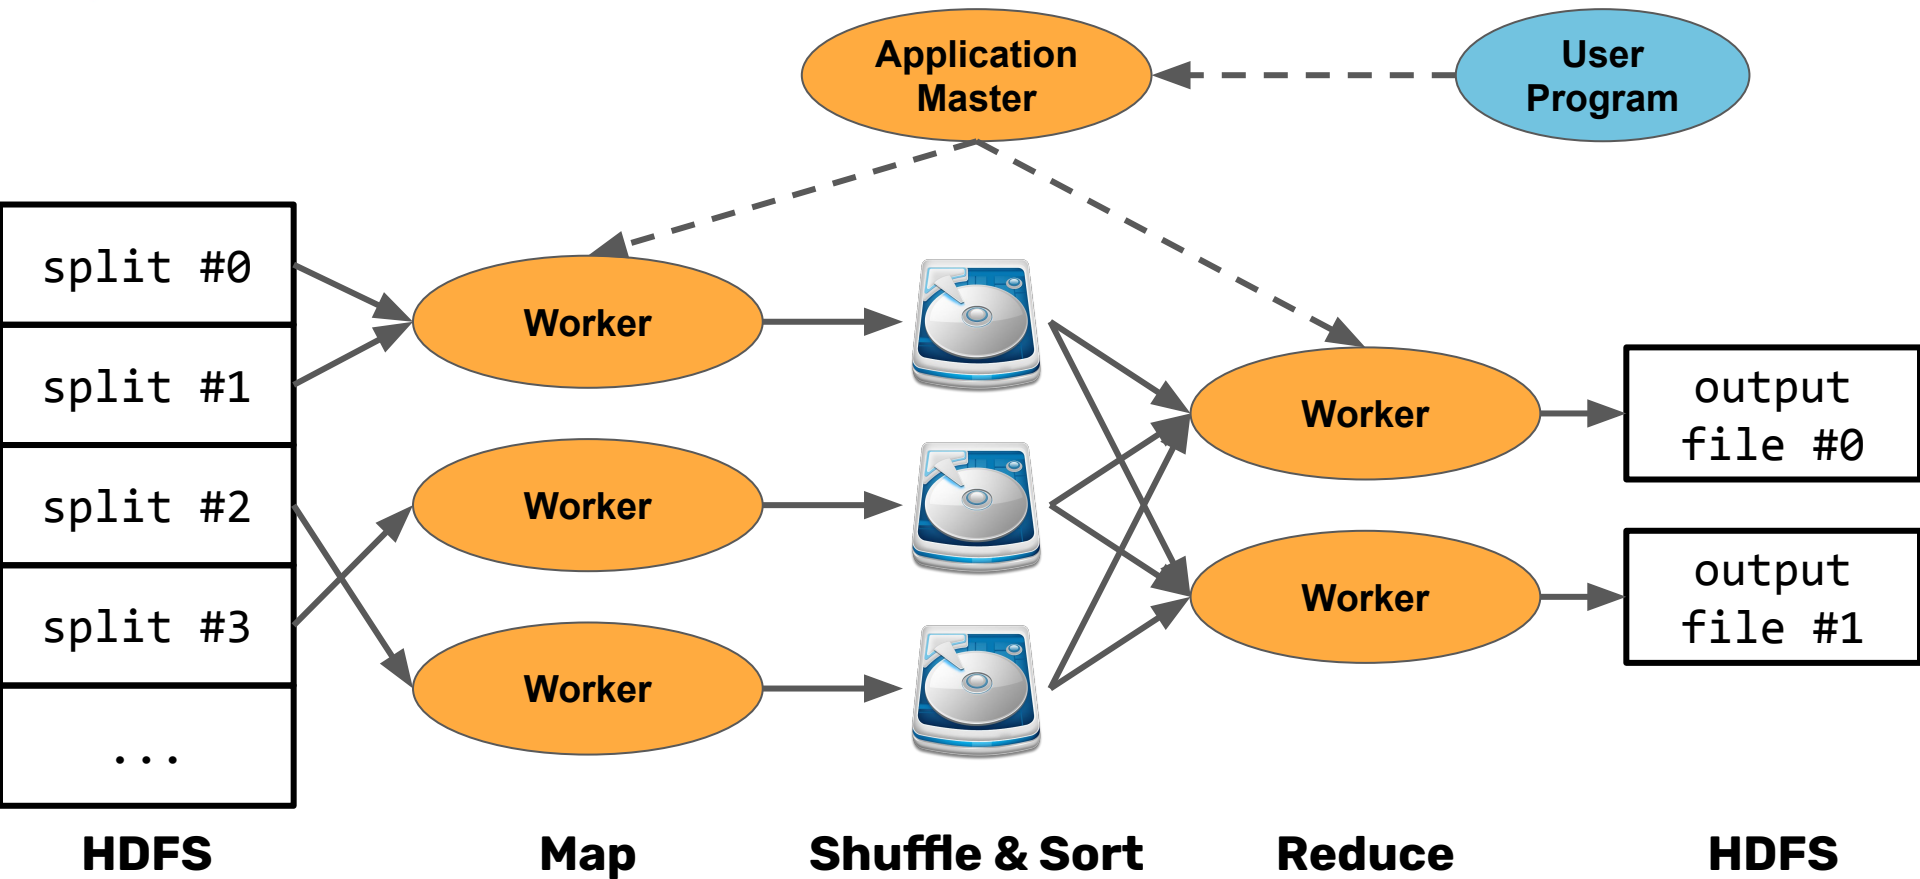
Reduce (фаза)



Fault Tolerance

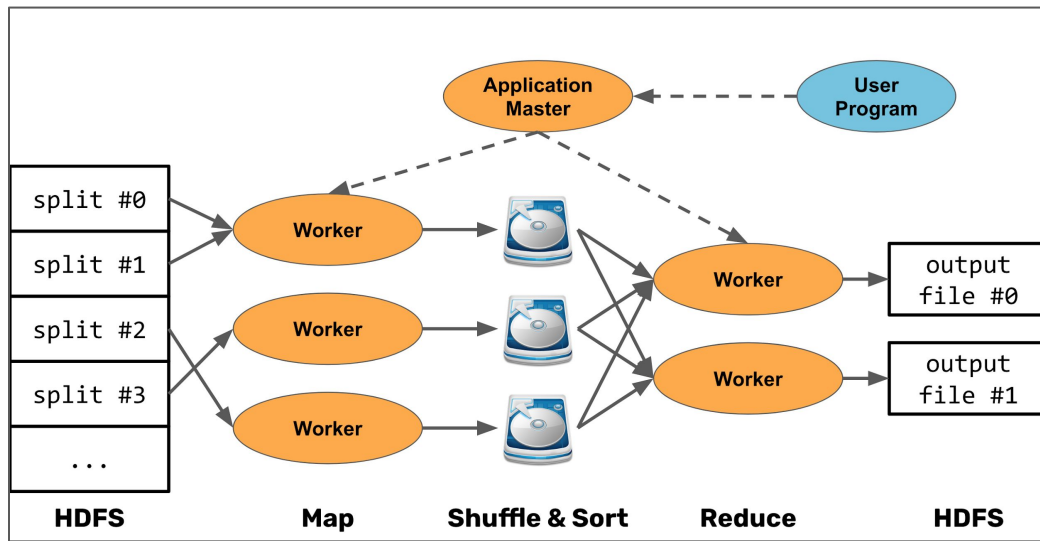


MapReduce





Где храним Shuffle & Sort?



A. RAM

Б. HDFS

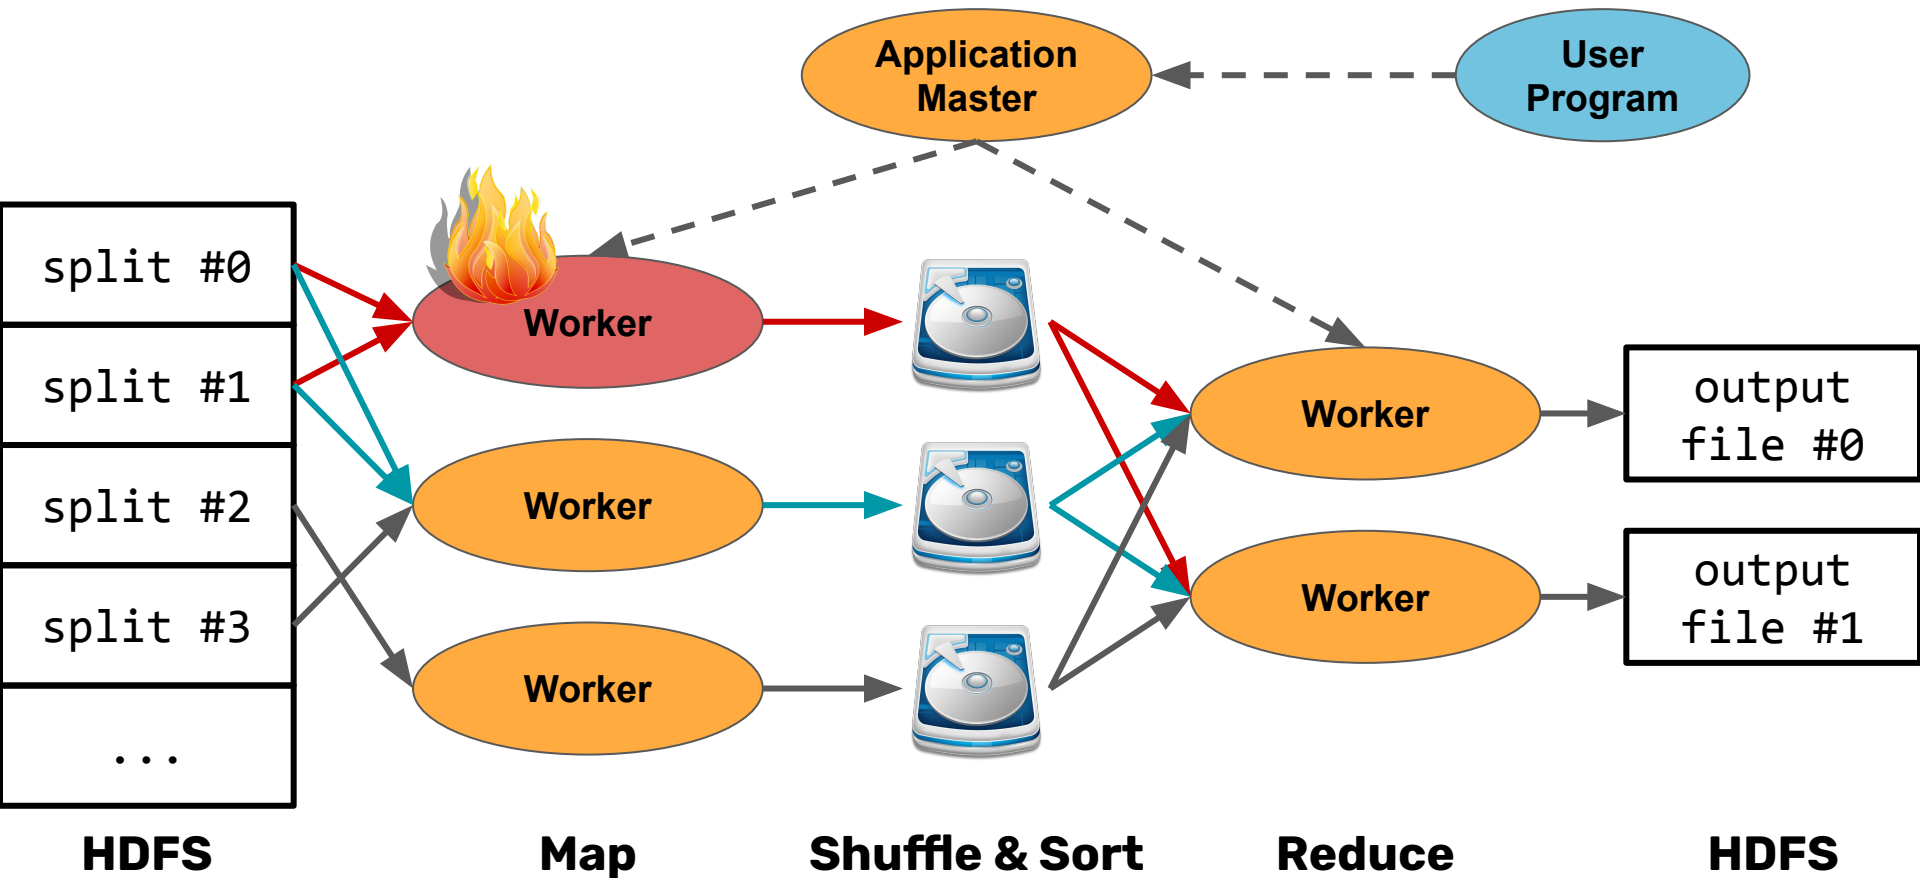
В. Local FS

Г. Где-то там



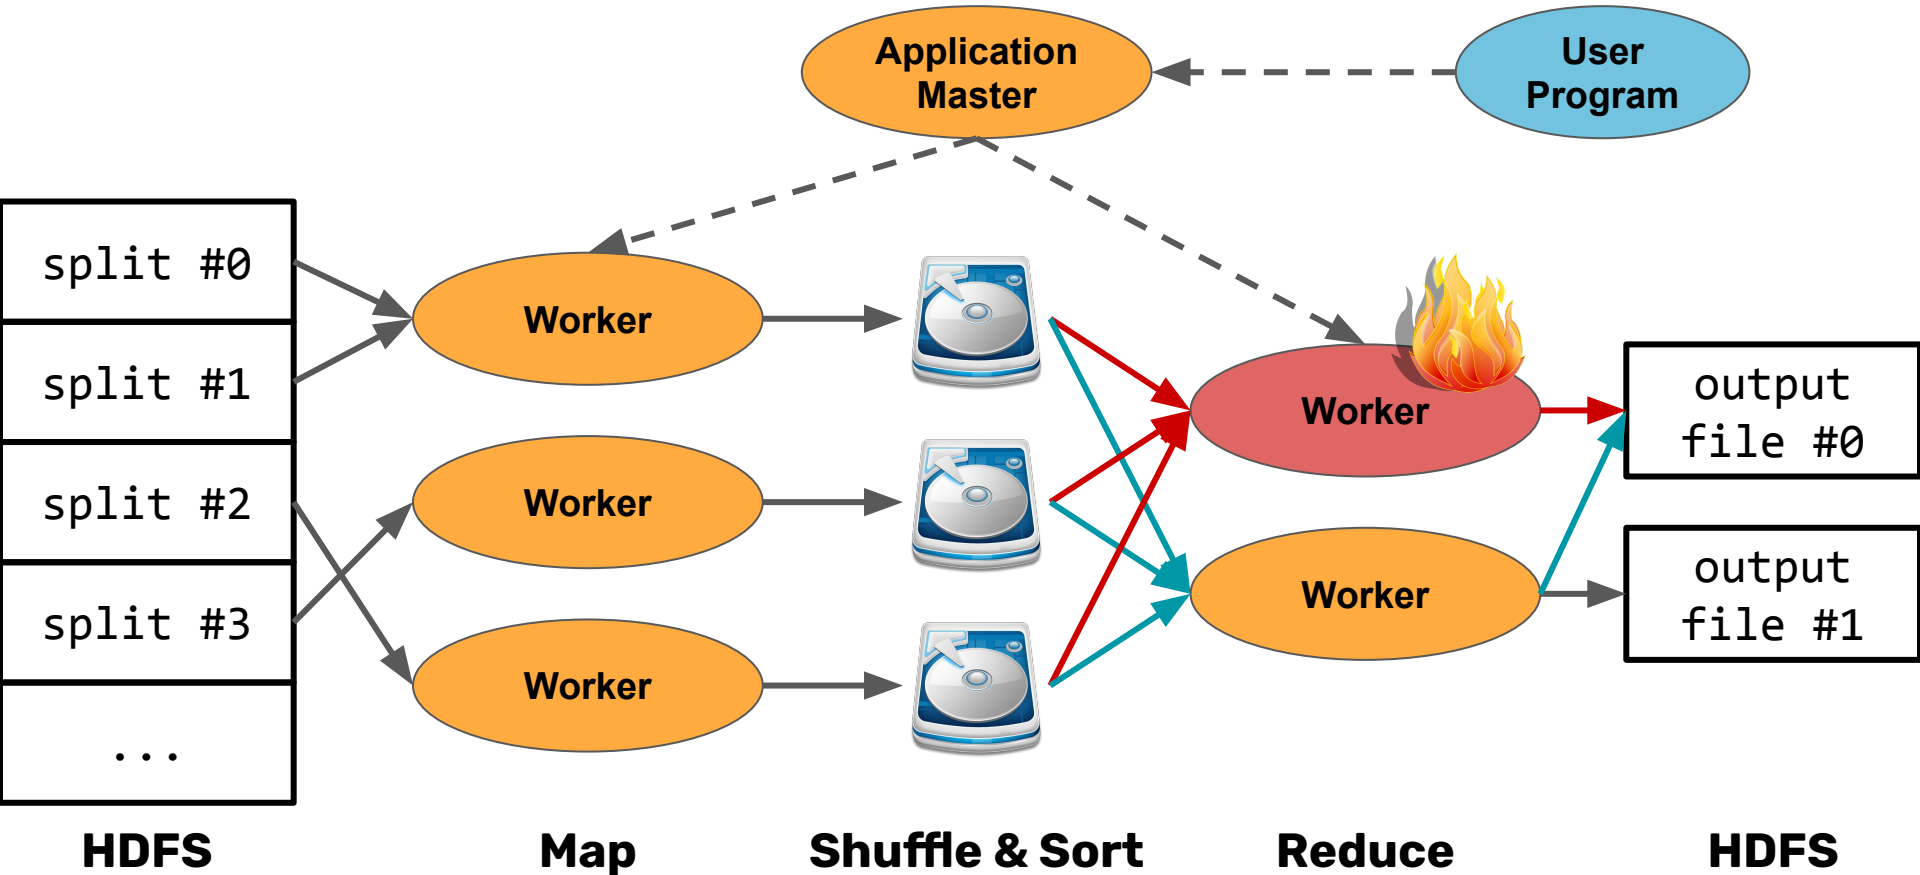


Fault Tolerance (защита от сбоев)



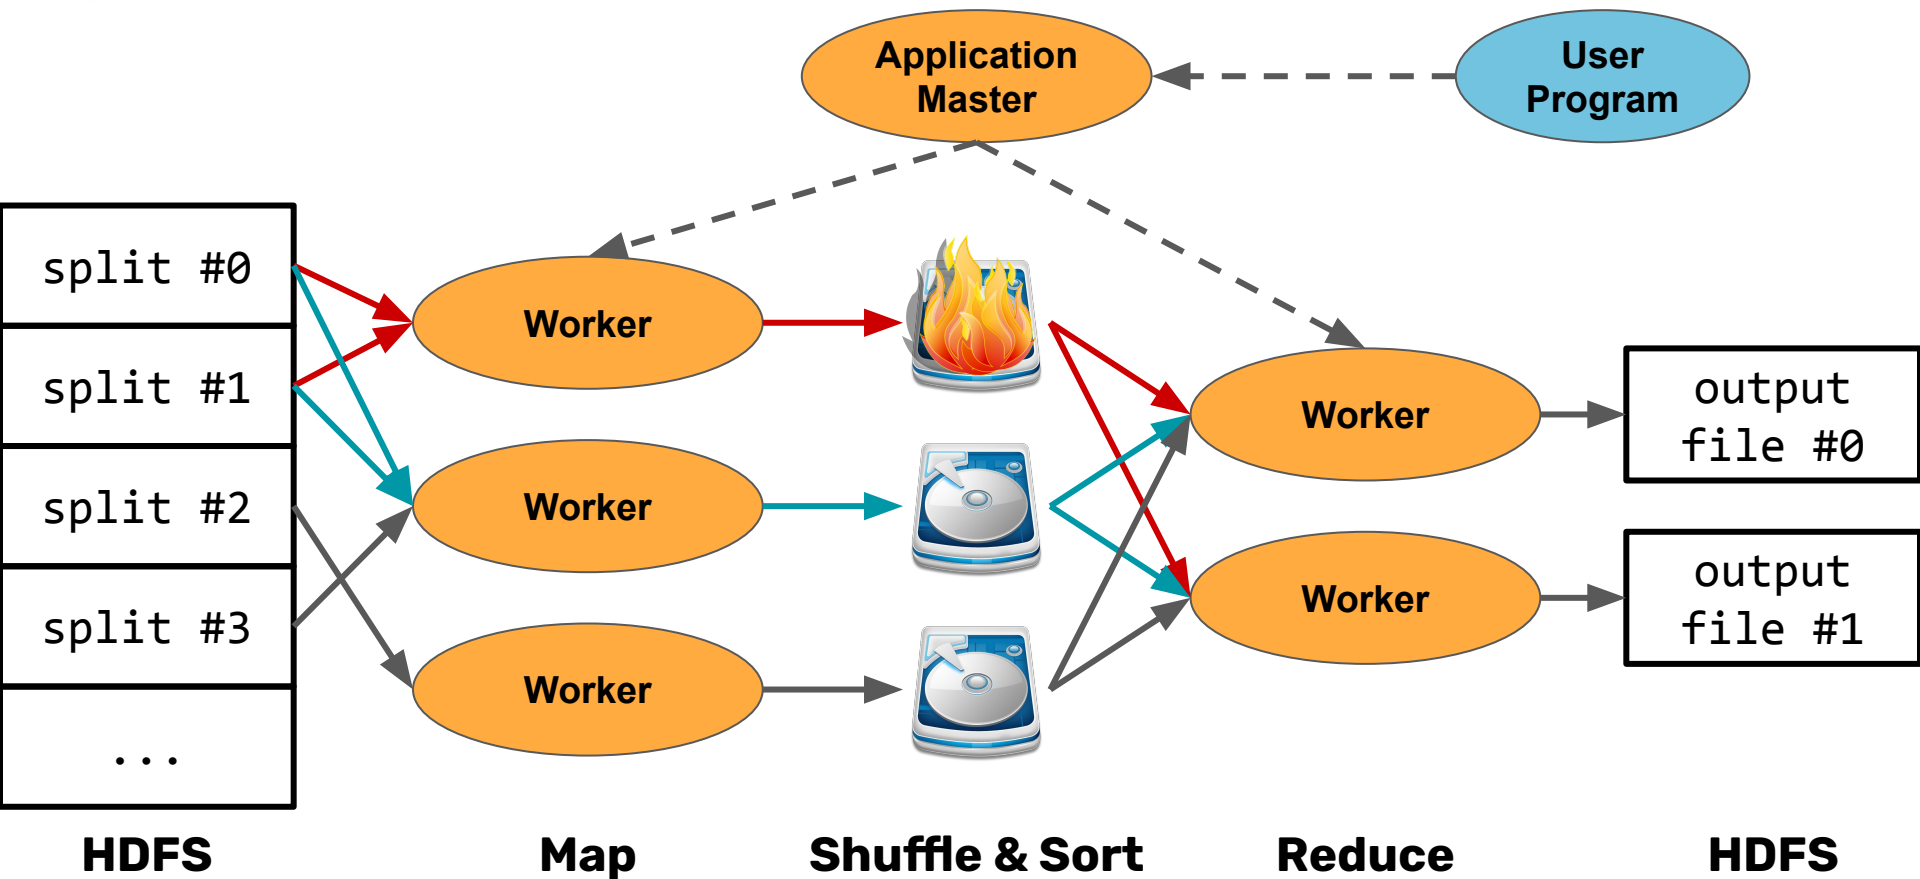


Fault Tolerance (защита от сбоев)



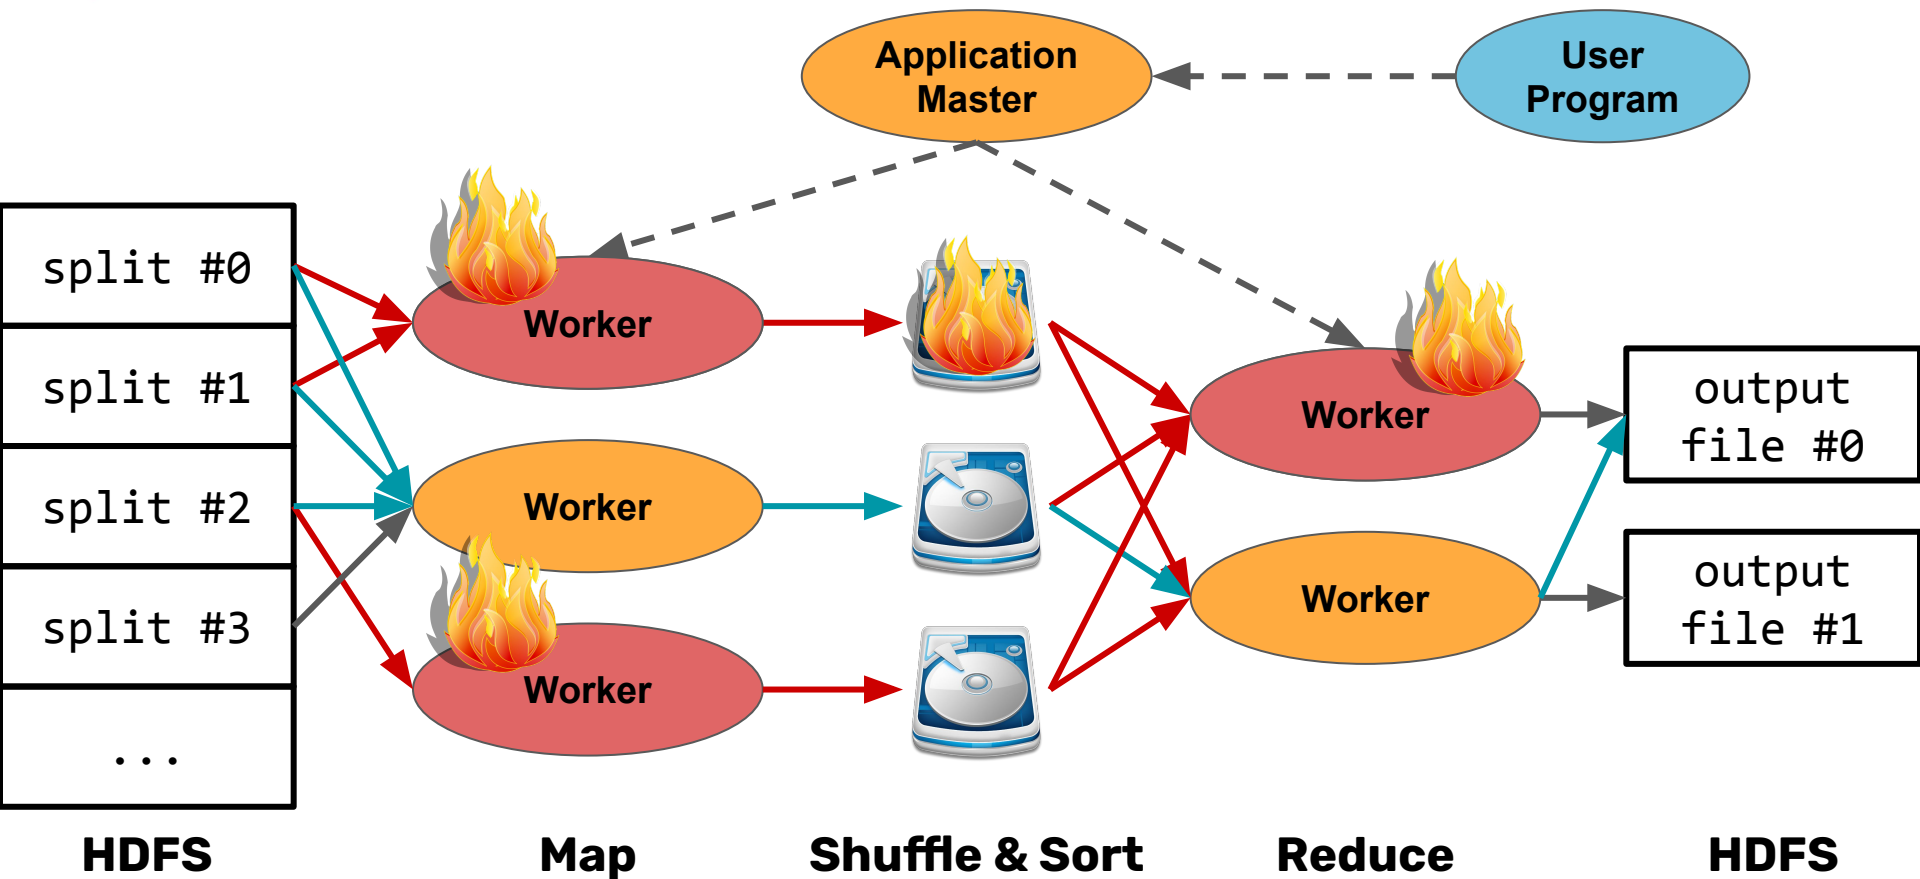


Fault Tolerance (защита от сбоев)





Fault Tolerance (защита от сбоев)



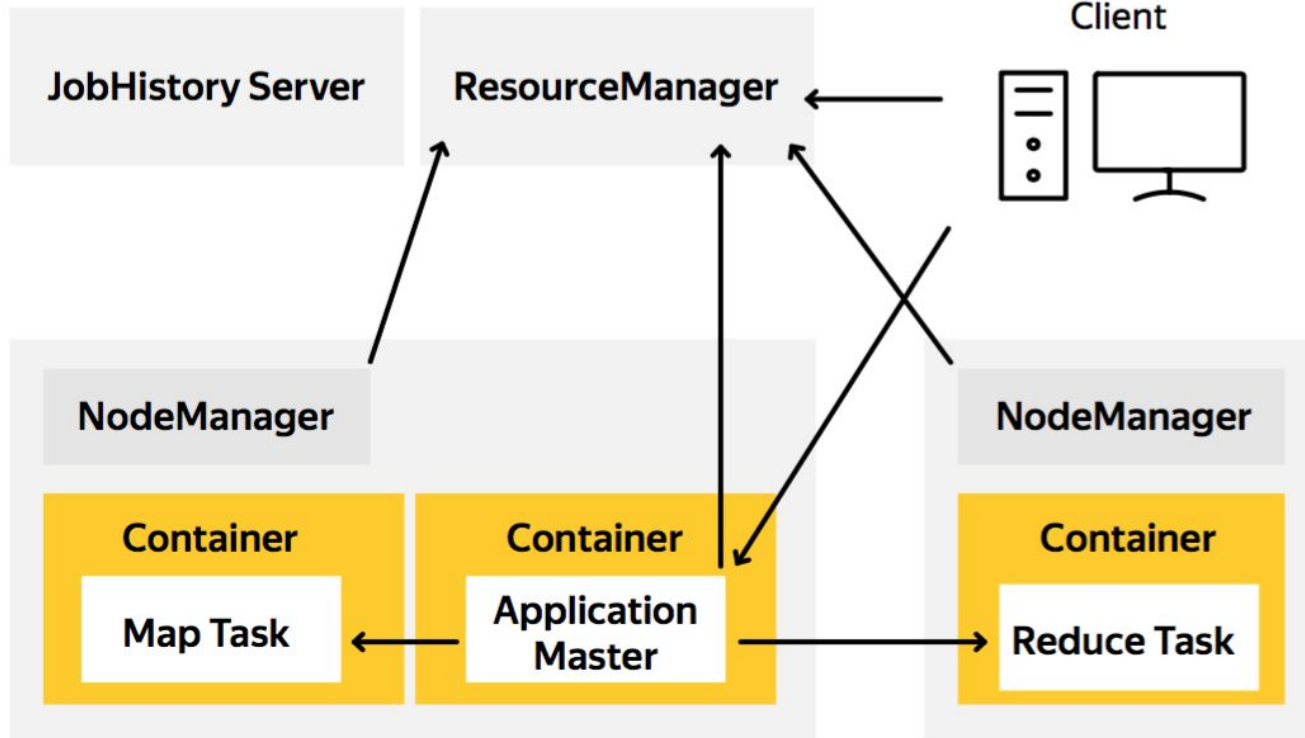


BIGDATA

**YOUR MAPREDUCE JOB WILL FINISH
SOMETIME**




THIS YEAR





MapReduce Frameworks (Timeline)

- ▶ [2003] Google MapReduce (development)
- ▶ [2004] Google MapReduce (article)
- ▶ [2005] Yandex MapReduce (development)
- ▶ [2006] Hadoop 0.1.0 release
- ▶ [2009] Amazon EMR (Hadoop inside)
- ▶ [2012] MapReduce → YARN
- ▶ [2013] Yahoo! YARN deployed in production
- ▶ [2013] Yandex YT...
- ▶ MapReduce in MongoDB, Riak, ...





Q&A

Как зная топологию данных
оптимизировать MapReduce?



```
$ yarn jar map_reduce_example.jar
```

```
...
```

```
INFO mapreduce.Job: Counters: 30
```

```
...
```

```
  Job Counters
```

```
    Launched map tasks=2
```

```
    Data-local map tasks=2
```

```
    Total time spent by all maps in occupied slots (ms)=27360
```

```
...
```



Tea / Coffee Break





MapReduce Streaming



MapReduce Streaming


input stream of key-value pairs





MapReduce Streaming

map: (k_in, v_in) --> [(k_interm, v_interm), ...]





MapReduce Streaming

aggregate by key (Shuffle & Sort)



<key,value>

Mapper

<key,[value1,value2...]>

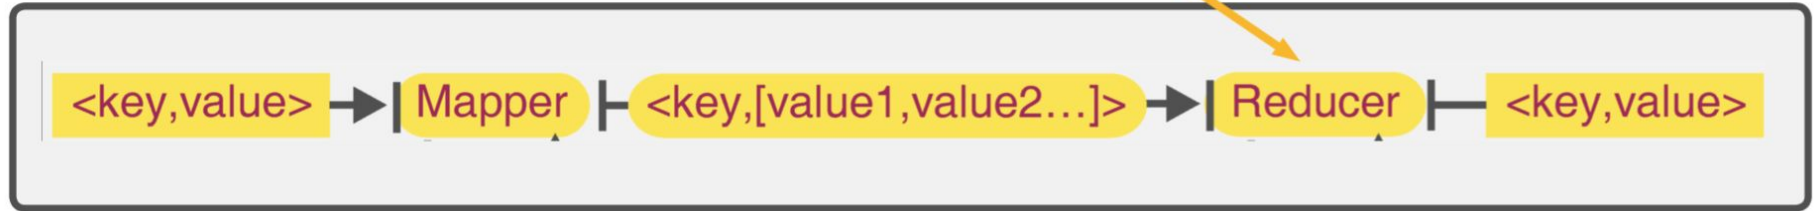
Reducer

<key,value>



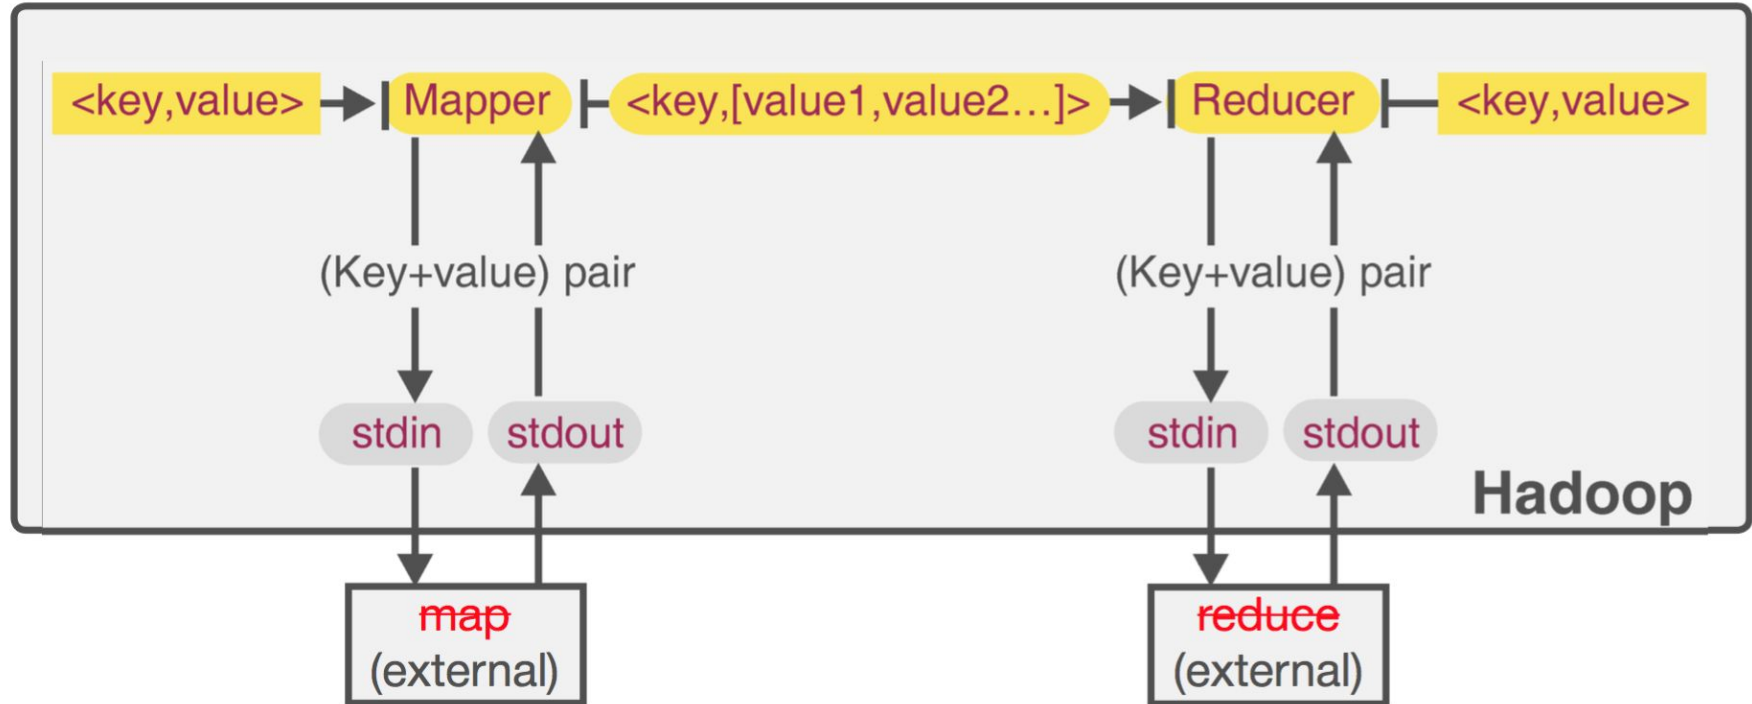
MapReduce Streaming

`reduce: (k_interm, [(v_interm, ...)]) --> [(k_out, v_out), ...]`



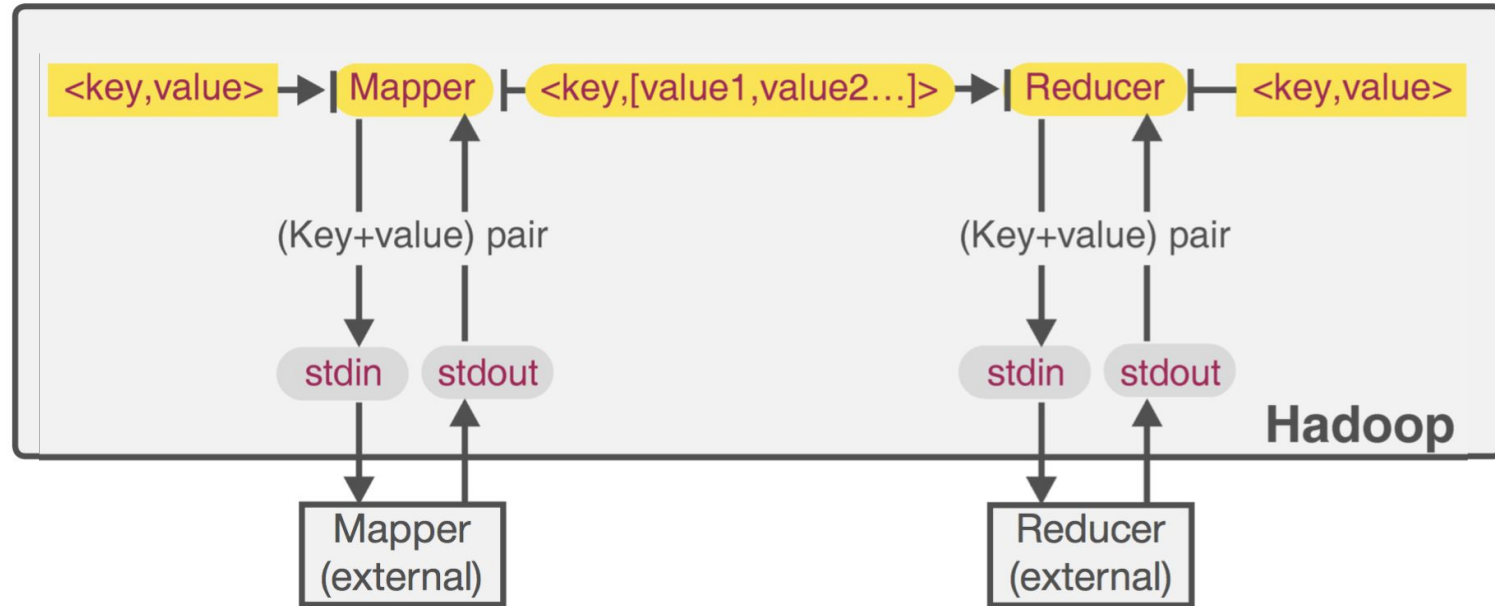


MapReduce Streaming





MapReduce Streaming



Mapper:

- ▶ Как данные читаем (input format)
- ▶ Как данные обрабатываем
- ▶ Как данные выводим (output format)

Тоже, что и Mapper, плюс:

- ▶ Как агрегируем по ключам отсортированные данные



**BIGDATA
TEAM**

Постановка задачи Line Count



WIKIPEDIA

`<article_id> <tab> <article_content>`



(k_{in}, v_{in})

map (функция)

```
map =  
(-, line) → [(-, 1), ...]
```

$[(k_{interm}, v_{interm}), ...]$

Map (фаза)

sort and group by
 k_{interm}

Shuffle & Sort

$(k_{interm}, [v_{interm}, ...])$

reduce (функция)

```
reduce = (uniq -c)  
(-, [1,1,...]) → (256, -)
```

$[(k_{out}, v_{out}), ...]$

Reduce (фаза)



```
run.sh
```

```
HADOOP_STREAMING_JAR=/path/to/hadoop-streaming.jar
```

```
OUT_DIR=my_hdfs_output
```

```
yarn jar $HADOOP_STREAMING_JAR \  
  -mapper "wc -l" \  
  -numReduceTasks 0 \  
  -input /data/wiki/en_articles_part \  
  -output $OUT_DIR
```




```
$ hdfs dfs -ls my_hdfs_output
```

```
Found 3 items
```

```
-rw-r--r--    3 aadral hdfs    0 2021-02-15 18:38 my_hdfs_output/_SUCCESS
-rw-r--r--    3 aadral hdfs    6 2021-02-15 18:38 my_hdfs_output/part-00000
-rw-r--r--    3 aadral hdfs    5 2021-02-15 18:38 my_hdfs_output/part-00001
```

```
$ hdfs dfs -text my_hdfs_output/*
```

```
3624
```

```
476
```



```
$ hdfs dfs -text my_hdfs_output/*
```

```
3624
```

```
476
```

```
$ hdfs dfs -ls -h /data/wiki/en_articles_part
```

```
Found 1 items
```

```
-rw-r--r--    3 hdfs hdfs      73.3 M  2020-03-12  21:03  
/data/wiki/en_articles_part/articles-part
```



```
$ ./run.sh
```

```
...
```

```
ERROR streaming.StreamJob: Error Launching job : Output directory  
hdfs://brain-master.bigdatateam.org:8020/user/aadral/my_hdfs_output  
already exists
```

```
Streaming Command Failed!
```



```
run.sh
```

```
HADOOP_STREAMING_JAR=/path/to/hadoop-streaming.jar
```

```
OUT_DIR=my_hdfs_output
```

```
hdfs dfs -rm -r $OUT_DIR
```

```
yarn jar $HADOOP_STREAMING_JAR \
```

```
-mapper "wc -l" \
```

```
-numReduceTasks 0 \
```

```
-input /data/wiki/en_articles_part \
```

```
-output $OUT_DIR
```



```
run.sh
```

```
HADOOP_STREAMING_JAR=/path/to/hadoop-streaming.jar
```

```
OUT_DIR=my_hdfs_output
```

```
hdfs dfs -rm -r $OUT_DIR
```

```
yarn jar $HADOOP_STREAMING_JAR \  
  -mapper "wc -l" \  
  -reducer "awk '{line_count += \$1} END { print line_count }'" \  
  -numReduceTasks 1 \  
  -input /data/wiki/en_articles_part \  
  -output $OUT_DIR
```



```
$ hdfs dfs -ls my_hdfs_output
```

Found 2 items

-rw-r--r--	3	aadral	hdfs	0	2021-02-17 11:22	my_hdfs_output/_SUCCESS
-rw-r--r--	3	aadral	hdfs	6	2021-02-17 11:22	my_hdfs_output/part-00000

```
$ hdfs dfs -text my_hdfs_output/*
```

```
4100
```



```
reducer.sh
```

```
#!/usr/bin/env bash
```

```
awk '{line_count += $1} END { print line_count }'
```



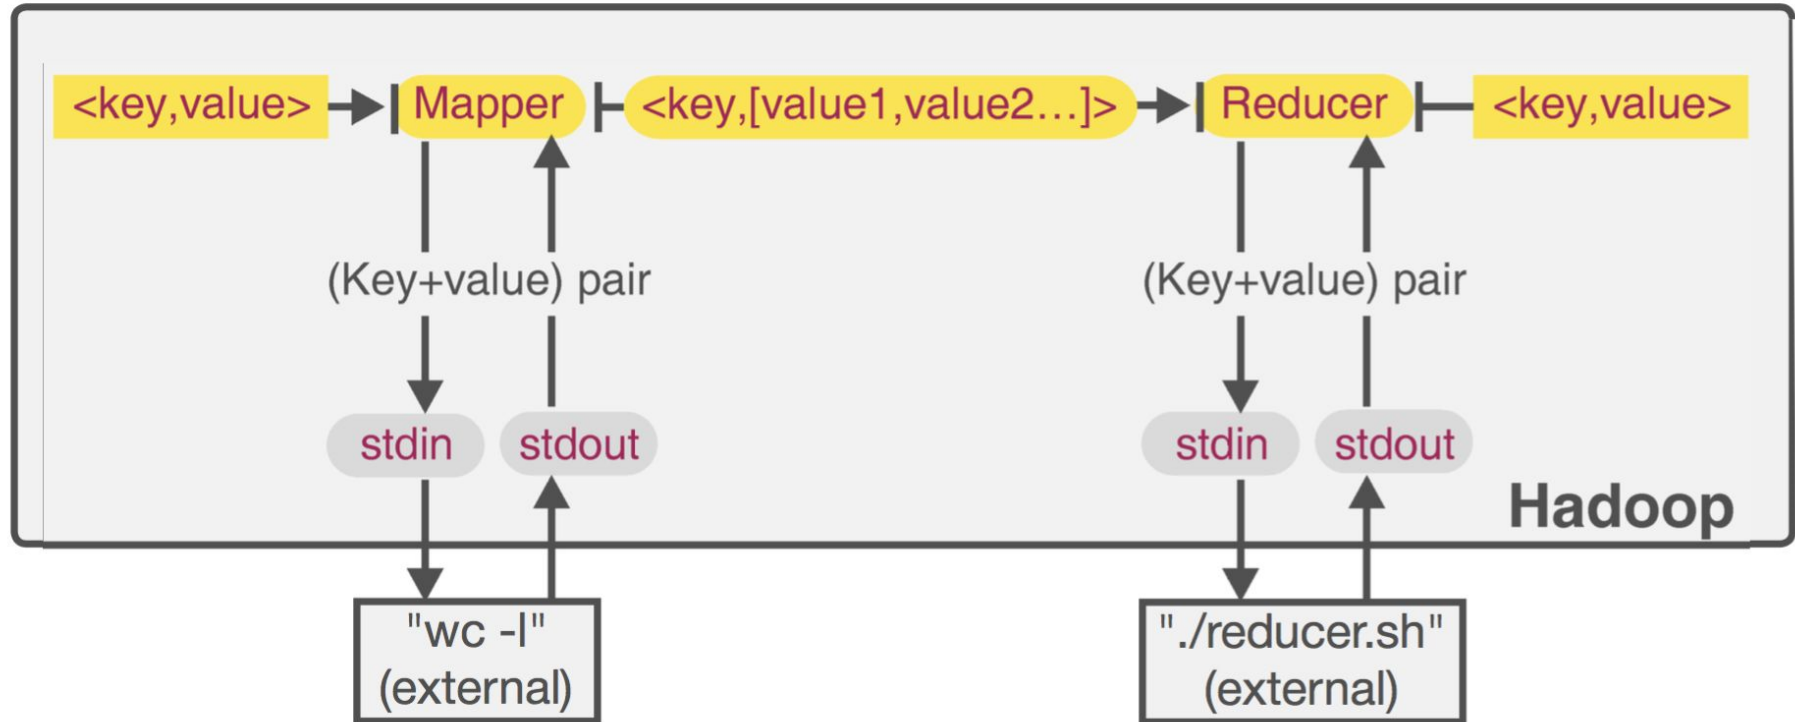
```
run.sh
```

```
HADOOP_STREAMING_JAR=/path/to/hadoop-streaming.jar
```

```
OUT_DIR=my_hdfs_output
```

```
hdfs dfs -rm -r $OUT_DIR
```

```
yarn jar $HADOOP_STREAMING_JAR \  
-files reducer.sh \  
-mapper "wc -l" \  
-reducer "./reducer.sh" \  
-numReduceTasks 1 \  
-input /data/wiki/en_articles_part \  
-output $OUT_DIR
```



- ▶ Вы можете объяснить, что происходит когда “умирает” Mapper или Reducer
- ▶ Вы знаете, за что отвечают ResourceManager и NodeManager в YARN
- ▶ Вы знаете 3 фазы MapReduce (Map, Shuffle & Sort, Reduce)
- ▶ Вы знаете, что такое MapReduce Streaming и как он работает (пример: Line Count)



KEEP
CALM
AND
TRY
CODING



**BIGDATA
TEAM**

Спасибо! Вопросы?

Feedback: http://rebrand.ly/x5bd2021q1_feedback_02_mr

Алексей Драль, aadral@bigdatateam.org

CEO at BigData Team, <http://bigdatateam.org>

<https://www.linkedin.com/in/alexey-dral>

<https://www.facebook.com/bigdatateam>