



RT Intro & Kafka

Vybornov Artyom, avybornov@bigdatateam.org

Big Data Instructor, <http://bigdatateam.org/>

Head of Data Platform, Rambler Group

<https://www.linkedin.com/in/artvybor/>





RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



Kafka

- ▷ Internal
- ▷ CLI



RT Intro

- ▷ **Batch -> RT**
- ▷ Approaches to RT data processing

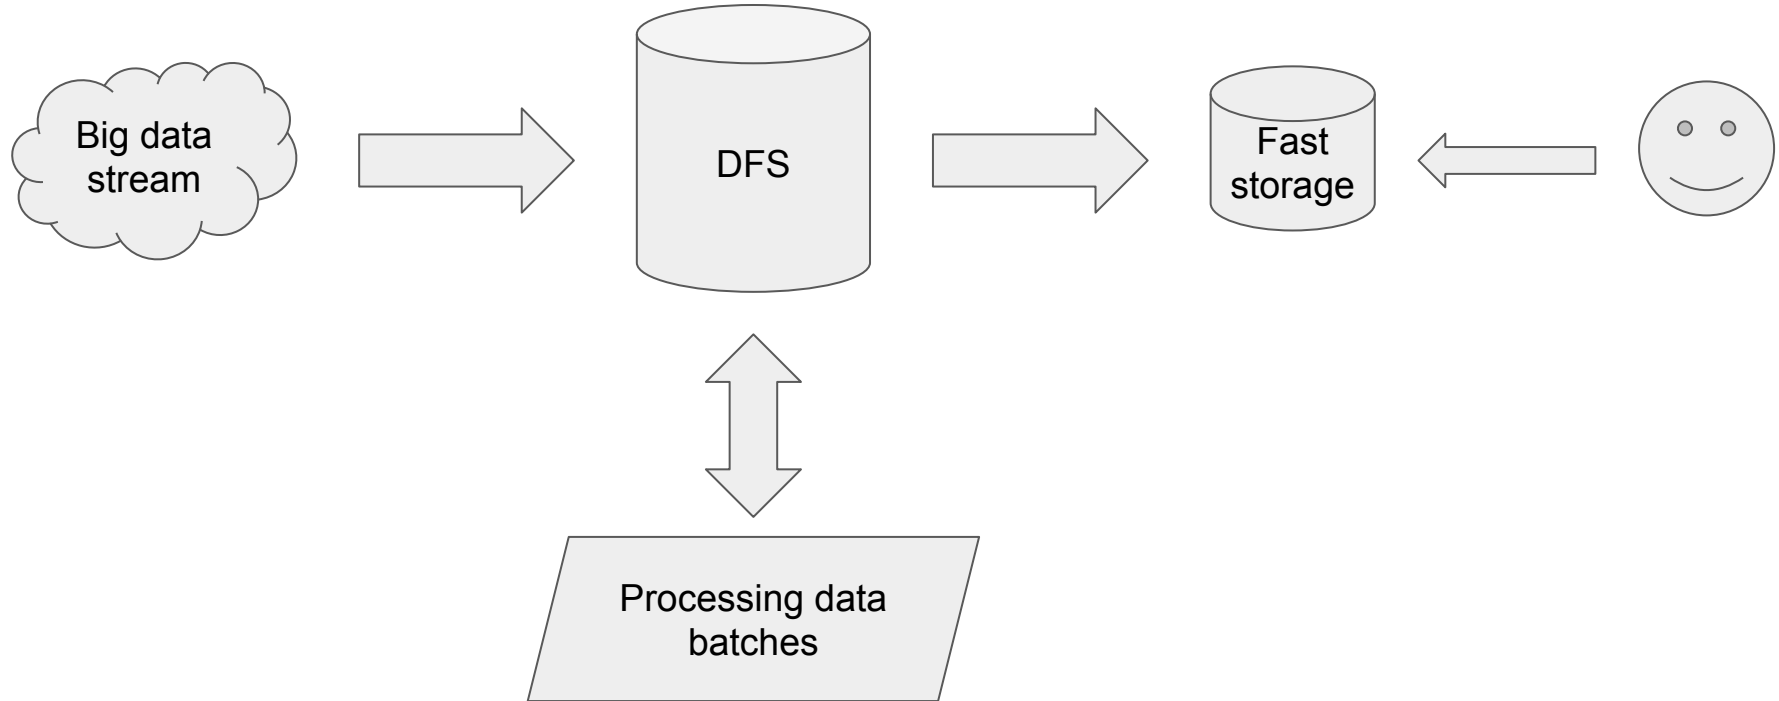


Kafka

- ▷ Internal
- ▷ CLI



Batch approach






The main cons of batch approach

- ▶ In practice a batch is a big time interval like an hour or day
- ▶ The size of batch is a minimal value for a lag
- ▶ Lag (delay) is the time between the event and its recording in the results of work
- ▶ For many purposes the following rule works: less lag => more valuable data



Real-time big data

- ▶ Real-time big data - is a set of technologies to process big data with minimal lag
- ▶ Without DFS
- ▶ Work with a stream of data instead of a batch





Real-time big data - lag in minutes



Build custom recommendations (Linkedin, Facebook)

- ▷ Billions of events per day
- ▷ Millions of events every second
- ▷ Minutes to build fresh recommendation



Real-time big data - lag in seconds



Programmatic advertising (Google, Facebook)

- ▶ Billions of events per day
- ▶ You liked the article and saw the relevant ads already on the next page



Credit card real-time fraud detection

- ▶ It takes a few seconds to detect a card with suspicious activity and block it



Real-time big data - lag in ms



Calls billing (mobile network operator)

- ▶ Hundred millions of users
- ▶ Less than a second to debit funds
- ▶ Zero fault tolerance



High-frequency trading (HFT)

- ▶ The reaction to the change of quotations of the exchange with a delay of 10-100 milliseconds



RT Intro

- ▷ Batch -> RT
- ▷ **Approaches to RT data processing**

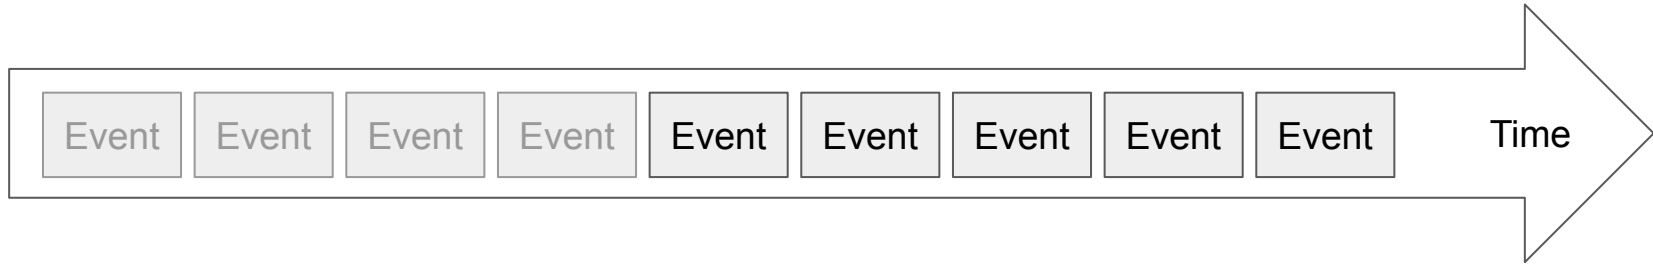


Kafka

- ▷ Internal
- ▷ CLI

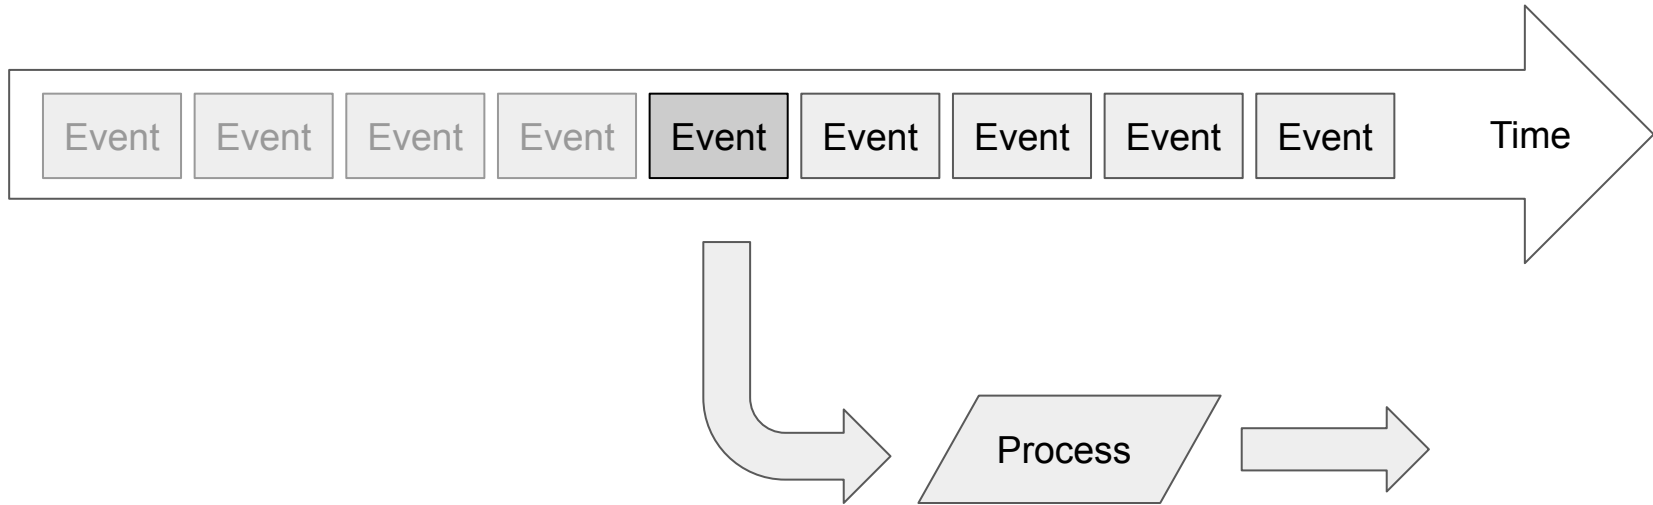


Event-based approach





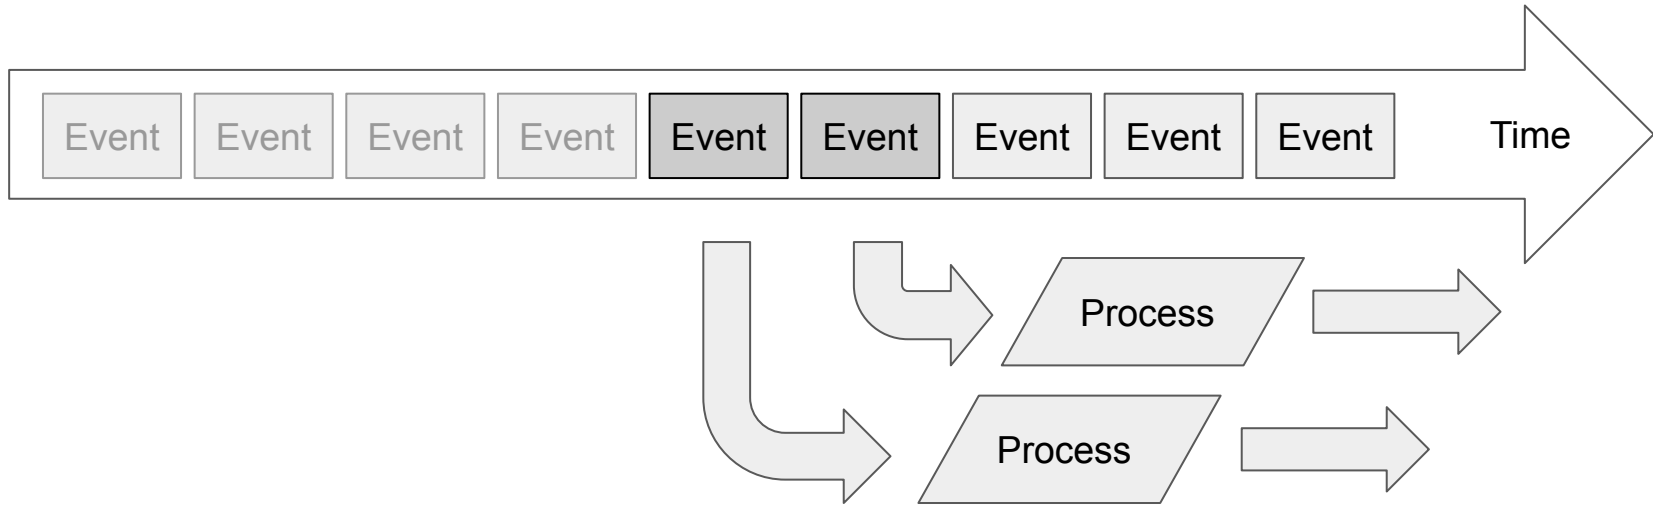
Event-based approach



Handling one event at a time




Event-based approach



- ▶ Handling one event at a time
- ▶ Events are processed in parallel, but completely independently
- ▶ Latency ~10ms



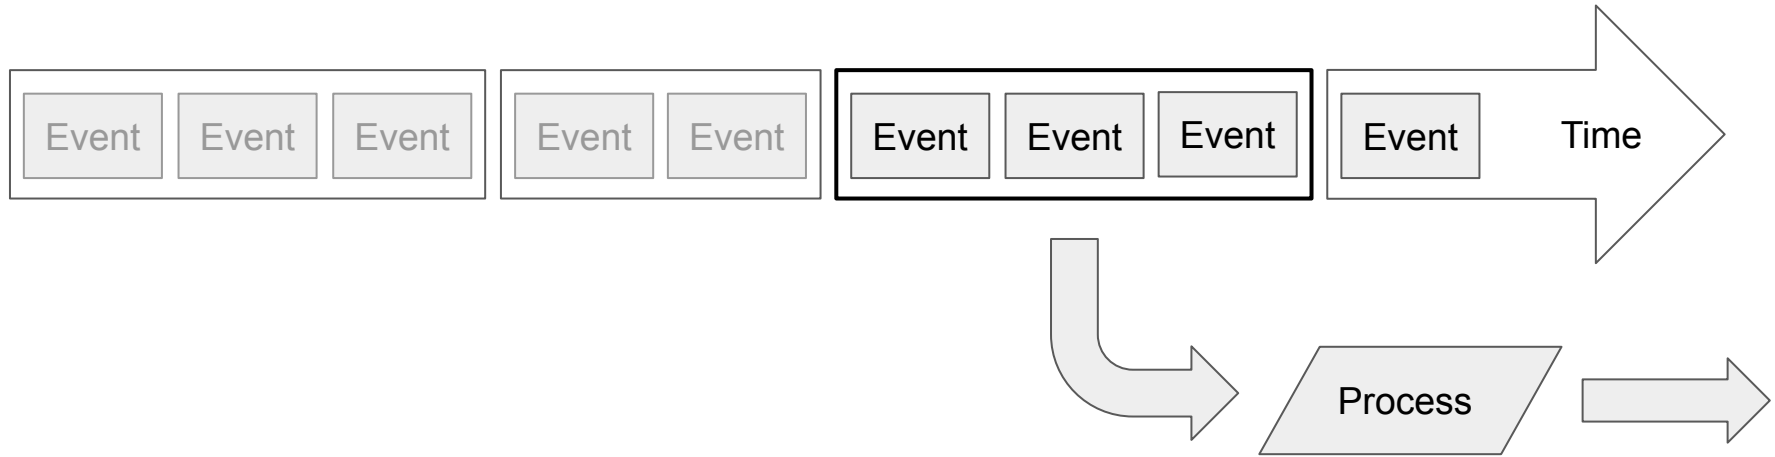
Micro-batch approach



- ▶ Stream is cut to batch by time (for example 10 seconds batch)



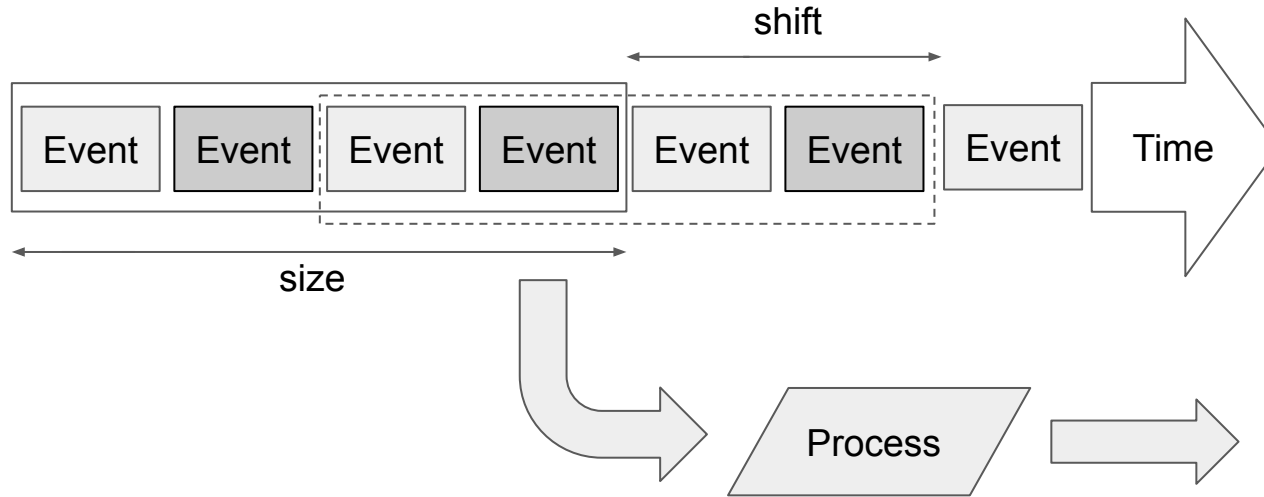
Micro-batch approach



- ▶ Stream is cut to batch by time (for example 10 seconds batch)
- ▶ Batches are processed sequentially
- ▶ Latency $\gg 1s$



Windowed approach



Batch as a sliding window



Similar to micro-batch approach





Event-based vs micro-batch

- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing



Event-based vs micro-batch

- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





Event-based vs micro-batch


- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





Event-based vs micro-batch


- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





Event-based vs micro-batch


- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





Event-based vs micro-batch


- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





Event-based vs micro-batch



- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





Event-based vs micro-batch


- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





Event-based vs micro-batch



- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





Event-based vs micro-batch


- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing





Event-based vs micro-batch

- ▶ Event-based allows you to achieve less lag
- ▶ Micro-batch allows you to save resources by reducing the common parts of each event handling/processing
- ▶ Micro-batch allows you to process more data on the same hardware






Throughput vs latency

- ▶ In real world resources are restricted
- ▶ Big data needs a huge throughput => in most cases we are choosing micro-batch
- ▶ There is no right answer - you should choose the approach by task



Throughput vs latency





RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing




Kafka

- ▷ Internal
- ▷ CLI



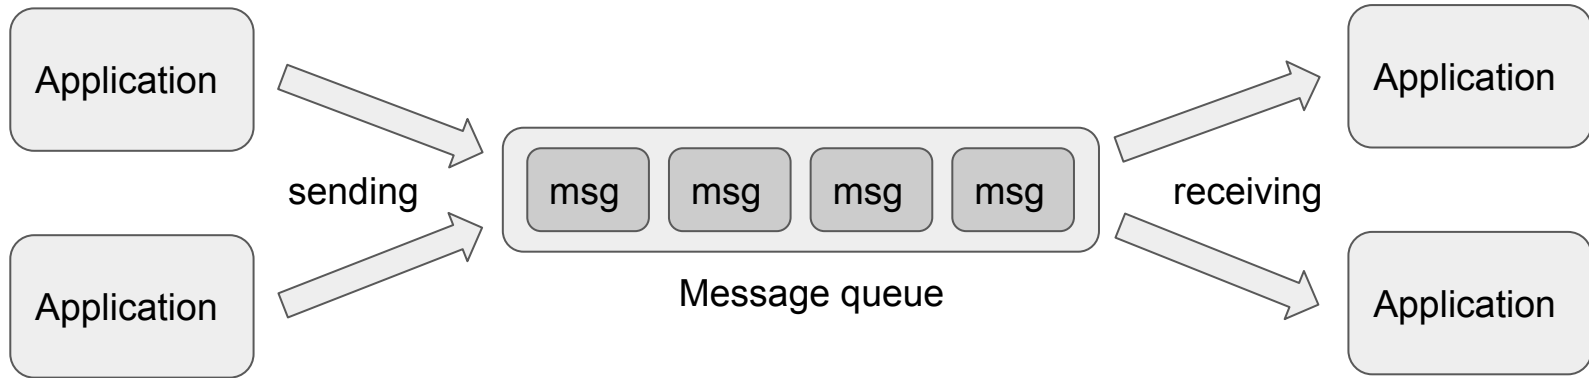
Requirements of storage for input data



- ▶ Event stream
- ▶ Big throughput (hundreds of thousands message per second)
- ▶ Small latency (less than 1s)



Stream handling - message queue



Message queue provides an asynchronous communications protocol between applications or between processes/threads inside a single application



Key features of classical message queue

- ▶ Complex schemes of message delivery
 - ▶ Reduces the throughput
- ▶ Per-message state
 - ▶ Reduces the throughput
- ▶ Stores the data in RAM
 - ▶ Not persistent storage
 - ▶ Strongly limits the amount of stored data



Storage of events in big data world

- ▶ Simplify the message delivery scheme
- ▶ Per message state not tracked by queue
- ▶ High throughput
- ▶ Persistent storage for a big amount of data




Google Cloud Pub/Sub





- ▶ Kafka is a unified, high-throughput, low-latency platform for handling real-time data feeds
- ▶ Kafka is a data bus for big data
- ▶ Kafka is an input events storage for real-time processing
- ▶ Kafka is an event-based real-time processing engine (Kafka Streams)







RT Intro & Kafka



RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing




Kafka

- ▷ **Internal**
- ▷ CLI

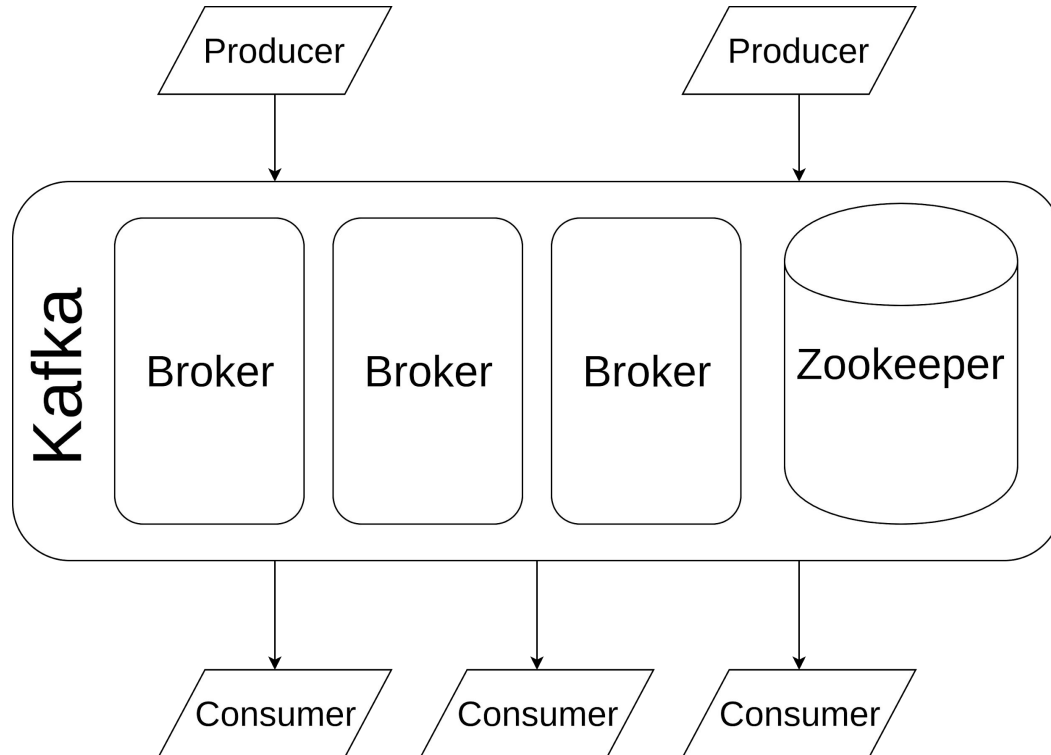


Kafka architecture





Kafka architecture





Topic



Partition 0

0	1	2	3
---	---	---	---

Partition 1

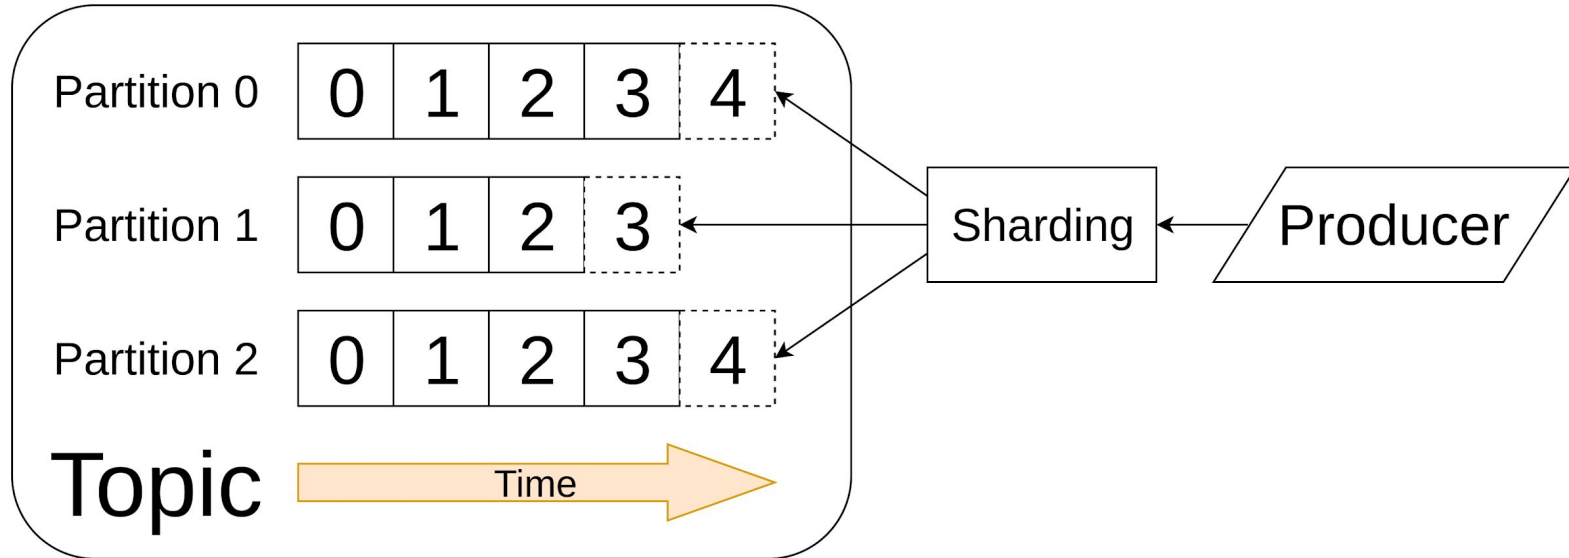
0	1	2
---	---	---

Partition 2

0	1	2	3
---	---	---	---

Topic




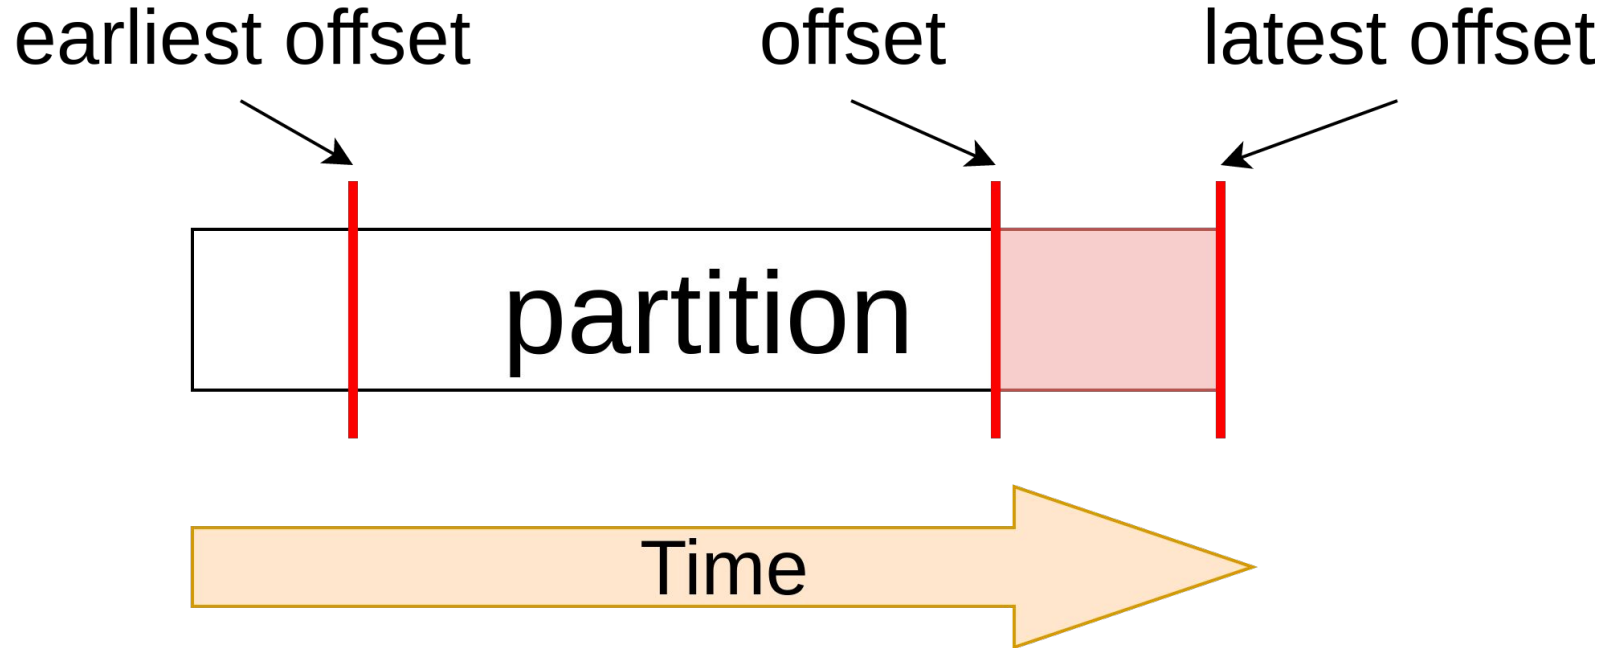






earliest offset


latest offset












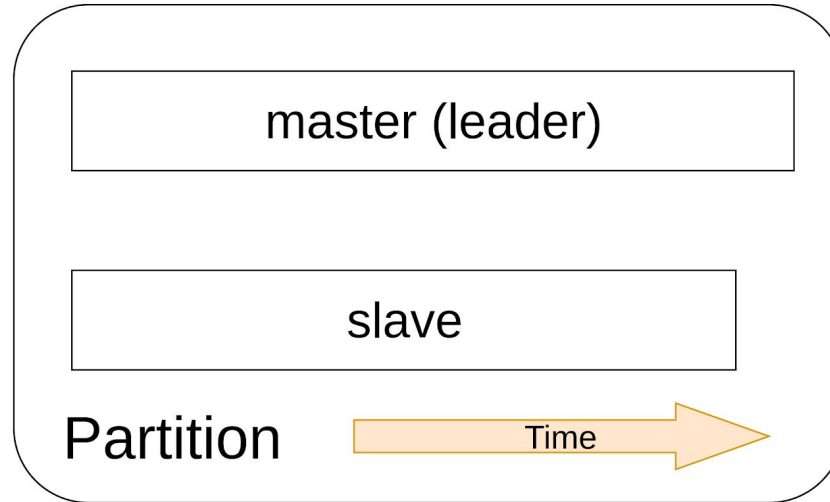


Kafka replication



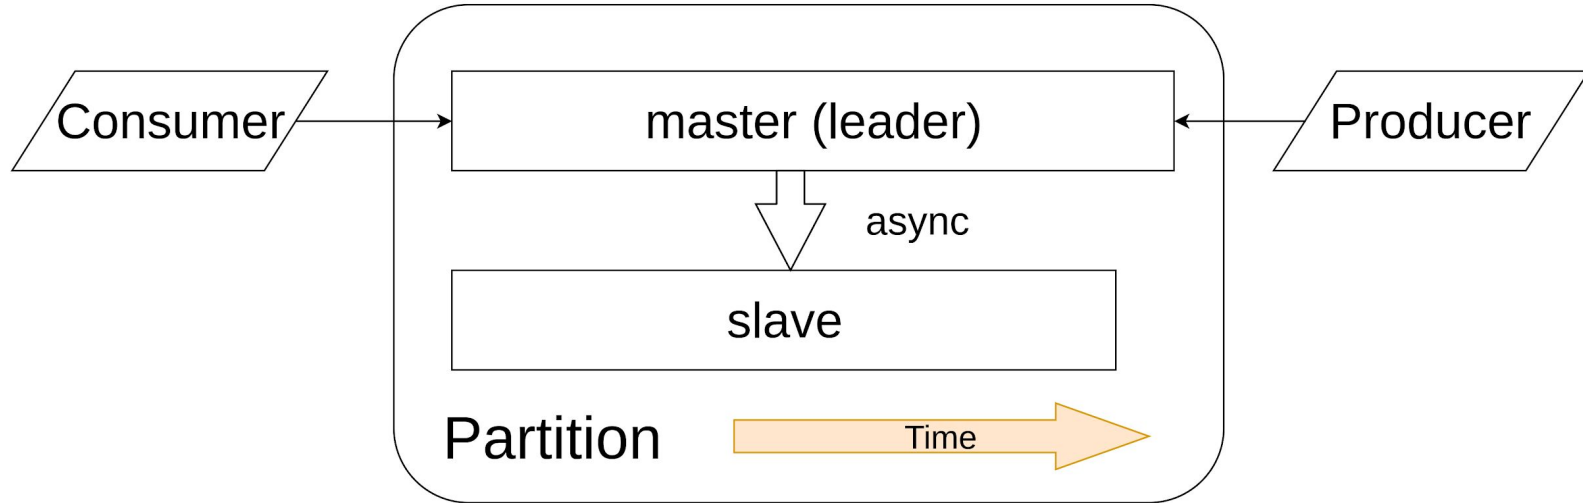


Kafka replication






Kafka replication





Kafka replication delay





RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



Kafka

- ▷ Internal
- ▷ **CLI**



Kafka CLI: kafka-topics

- ▶ It is the utility to create, delete, describe or change the topic
- ▶ `kafka-topics.sh --bootstrap-server $BROKERS --create --topic test_topic --partitions 3 --replication-factor 2`
- ▶ `kafka-topics.sh --bootstrap-server $BROKERS --describe --topic test_topic`
- ▶ `kafka-topics.sh --bootstrap-server $BROKERS --list`



Kafka CLI: kafka-console-producer



It is utility to send data from standard input and to Kafka topic



```
kafka-console-producer.sh --broker-list $BROKERS --topic  
test_topic
```



Kafka CLI: kafka-console-consumer



It is utility to read data from Kafka topic



```
kafka-console-consumer.sh --bootstrap-server $BROKERS  
--topic test_topic --from-beginning
```



Kafka CLI: kafka-run-class



It is entry point to run any class in the Kafka environment



```
kafka-run-class.sh kafka.tools.GetOffsetShell  
--broker-list $BROKERS --topic test_topic --time -1
```




RT Intro

- ▷ Batch -> RT
- ▷ Approaches to RT data processing



Kafka

- ▷ Internal
- ▷ CLI



Thank you! Questions?

Vybornov Artyom, avybornov@bigdatateam.org

Big Data Instructor, <http://bigdatateam.org/>

Head of Big Data Dev Team, Rambler Group

<https://www.linkedin.com/in/artvybor/>